OpenDiscoveryTrace: Process Traces for Evaluating AI Scientist Workflows

Anonymous Authors[1]

Abstract

We propose OPENDISCOVERYTRACE, a dataset of 522 AI scientific agent trajectories capturing *how* models reason—not just what they produce. Each trajectory records 9 fields per step across 124 tasks in four science domains, executed by 7 models (3 frontier, 4 open-weight). Pilot analysis shows identical success rates ($\sim$69%) yet 30$\times$ error divergence between models—a dimension invisible to output-only benchmarks. Released under CC-BY-4.0.

1. Scientific Bottleneck

AI systems now conduct scientific research at scale (Jumper et al., 2021; Merchant et al., 2023; Szymanski et al., 2023), and LLM-based agents orchestrate full research pipelines (Lu et al., 2024; Boiko et al., 2023). Benchmarks have emerged to evaluate these systems (Chen et al., 2025; Majumder et al., 2024; Starace et al., 2025; Huang et al., 2024), yet every one evaluates only *final outputs*—generated code, hypotheses, or rubric scores. The reasoning process is discarded.

This creates three problems. First, evaluating only outputs conflates reasoning quality with outcome luck. Second, governance requires complete decision traces for auditing (Kapoor & Narayanan, 2023; Hung & Yen, 2024; Tom et al., 2024). Third, improving agents requires understanding failure modes, which are trajectory properties, not output properties. OPENDISCOVERYTRACE addresses this gap with the first public dataset of structured AI scientific process traces.

2. Proposed Dataset

OPENDISCOVERYTRACE comprises **522 trajectories across 7 models**: 372 from three frontier models (GPT-5.4, Claude Opus 4.6, Gemini 3.1 Pro; 124 each, balanced across 4 domains $\times$ 3 difficulties), 120 from four open-weight models (Qwen2.5-7B, Mistral-7B-v0.3, Phi-3.5-mini, Qwen2.5-1.5B; 30 each), and 30 live-retrieval variants.

Each step records 9 fields extending ReAct (Yao et al., 2023): `phase`, `thought`, `action` (tool/input/output), `observation`, `error`, `revision_trigger`, `confidence`, `step_id`, `timestamp`. Five benchmark tasks are defined: outcome prediction, error localization, claim verification, autonomy classification, and process quality scoring. Full schema in Appendix A.

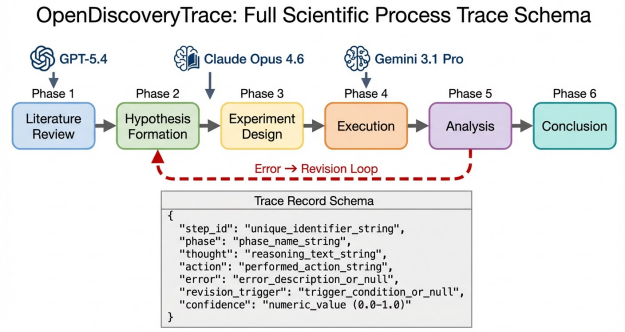

Figure 1. OPENDISCOVERYTRACE: 7 models execute tasks through a six-phase workflow. Each step is a 9-field trace. The 522 trajectories form a corpus of AI scientific process traces.

Why process traces matter. All three frontier models achieve indistinguishable success rates ($\sim$69%, p=0.997), yet Claude Opus 4.6 produces 30$\times$ more errors than GPT-5.4 (2.5 vs. 0.08; H=122, $p<10^{-4}$; Cliff's δ=0.613). These errors are qualitatively different: Claude's are 66.7% tool misuse; GPT-5.4's are 83.6% reasoning errors. Output-only benchmarks miss this entirely.

3. Acquisition Roadmap

Phase 1 (current). 522 trajectories, 7 models, 124 tasks, 4 domains. Cost: \$300 API + \$40 GPU. Schema, harness, benchmarks operational. **Phase 2 (months 1–2).** All 200 tasks $\times$ 9 models. Target: 1,800+ trajectories. **Phase 3 (months 2–3).** Human annotation on 120-trajectory sample (3 annotators, 5 axes; LLM pilot achieves α=0.63–0.82). **Phase 4 (months 3–6).** Live retrieval for literature tasks; pilot ELN integration (Dirnagl & Bhatt, 2016). **Sustain-**

[1]Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Submitted to the AI for Science workshop (ICML 2026).

ability. Semver-versioned schema on HuggingFace with DOI. Open-source harness. Annual model refresh.

4. Metadata, Governance & Safety

License: CC-BY-4.0. **Format:** Self-contained JSON per trajectory. Datasheet per Gebru et al. (2021). **Release tiers:** (1) Full traces where provider ToS permits; (2) abstracted summaries (phase transitions, tool calls, errors) preserving $>90\%$ analytical utility. **Privacy:** No PII; public scientific knowledge only. **Splits:** Task-level 80/20 train/test, domain-stratified, no leakage (Kapoor & Narayanan, 2023). Evaluation follows HELM (Liang et al., 2023).

5. Acceleration Potential

OpenDiscoveryTrace enables: (1) process-aware evaluation that distinguishes reasoning from luck (Wang et al., 2023); (2) early failure intervention (68.1% of first errors at step 0); (3) error-type-specific model improvement; (4) process-aware reward modeling for RL; (5) trustworthy AI scientist auditing (King et al., 2009); (6) reproducible open-weight benchmarking (120 trajectories from 4 models, no API keys required).

References

Boiko, D. A., MacKnight, R., Kline, B., and Gomes, G. Autonomous chemical research with large language models. *Nature*, 624(7992):570–578, 2023. doi: 10.1038/s41586-023-06792-0.

Chan, J. S., Chowdhury, N., Jaffe, O., Aung, J., Sherburn, D., Mays, E., Starace, G., Liu, K., Maksin, L., Patwardhan, T., Weng, L., and Madry, A. MLE-bench: Evaluating machine learning agents on machine learning engineering. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2410.07095.

Chen, Z., Chen, S., Ning, Y., Zhang, Q., Wang, B., Yu, B., Li, Y., Liao, Z., Wei, C., Lu, Z., Dey, V., Xue, M., Baker, F. N., Burns, B., Adu-Ampratwum, D., Huang, X., Ning, X., Gao, S., Su, Y., and Sun, H. ScienceAgentBench: Toward rigorous assessment of language agents for data-driven scientific discovery. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2410.05080.

Dirnagl, U. and Bhatt, D. Electronic lab notebooks: Forget about putting pen to paper – it's time to go digital. *Nature*, 537(7618):168–169, 2016. doi: 10.1038/nj7618-168a.

Gebru, T., Morgenstern, J., Vecchione, B., Vaughan, J. W., Wallach, H., Daumé III, H., and Crawford, K. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021. doi: 10.1145/3458723.

Huang, Q., Vora, J., Liang, P., and Leskovec, J. MLAgentBench: Evaluating language agents on machine learning experimentation. *arXiv preprint arXiv:2310.03302*, 2024.

Hung, Y. and Yen, C.-C. Levels of AI agents: from rules to large language models. *arXiv preprint arXiv:2405.06643*, 2024.

Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., Tunyasuvunakool, K., Bates, R., Žídek, A., Potapenko, A., et al. Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873):583–589, 2021. doi: 10.1038/s41586-021-03819-2.

Kapoor, S. and Narayanan, A. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns*, 4(9):100804, 2023. doi: 10.1016/j.patter.2023.100804.

King, R. D., Rowland, J., Oliver, S. G., Young, M., Aubrey, W., Byrne, E., Liakata, M., Markham, M., Pir, P., Soldatova, L. N., Sparkes, A., Whelan, K. E., and Clare, A. The automation of science. *Science*, 324(5923):85–89, 2009. doi: 10.1126/science.1165620.

Liang, P., Bommasani, R., Lee, T., Tsipras, D., Soylu, D., Yasunaga, M., Zhang, Y., Narayanan, D., Wu, Y., Kumar, A., Newman, B., Yuan, B., Yan, B., Zhang, C., Cosgrove, C., Manning, C. D., Ré, C., Acosta-Navas, D., Hudson, D. A., Zelikman, E., Durmus, E., Ladhak, F., Rong, F., Ren, H., Yao, H., Wang, J., Santhanam, K., Orber, L. J., Hallman, M., et al. Holistic evaluation of language models. *Transactions on Machine Learning Research*, 2023.

Liu, X., Yu, H., Zhang, H., Xu, Y., Lei, X., Lai, H., Gu, Y., Ding, H., Men, K., Yang, K., Zhang, S., Deng, X., Zeng, A., Du, Z., Zhang, C., Shen, S., Zhang, T., Su, Y., Sun, H., Huang, M., Dong, Y., and Tang, J. AgentBench: Evaluating LLMs as agents. *arXiv preprint arXiv:2308.03688*, 2024.

Liu, Y., Yang, Z., Xie, T., Ni, J., Gao, B., Li, Y., Tang, S., Ouyang, W., Cambria, E., and Zhou, D. ResearchBench: Benchmarking LLMs in scientific discovery via inspiration-based task decomposition. *arXiv preprint arXiv:2503.21248*, 2025.

Lu, C., Lu, C., Lange, R. T., Foerster, J., Clune, J., and Ha, D. The AI scientist: Towards fully automated open-ended scientific discovery. *arXiv preprint arXiv:2408.06292*, 2024.

Majumder, B. P., Surana, H., Agarwal, D., Mishra, B. D., Meena, A., Prakhar, A., Vora, T., Khot, T., Sabharwal, A., and Clark, P. DiscoveryBench: Towards data-driven discovery with large language models. *arXiv preprint arXiv:2407.01725*, 2024.

Merchant, A., Batzner, S., Schoenholz, S. S., Aykol, M., Cheon, G., and Cubuk, E. D. Scaling deep learning for materials discovery. *Nature*, 624(7990):80–85, 2023. doi: 10.1038/s41586-023-06735-9.

Starace, G., Jaffe, O., Sherburn, D., Aung, J., Chan, J. S., Maksin, L., Dias, R., Mays, E., Kinsella, B., Thompson, W., Heidecke, J., Glaese, A., and Patwardhan, T. PaperBench: Evaluating AI's ability to replicate AI research. *arXiv preprint arXiv:2504.01848*, 2025.

Szymanski, N. J., Rendy, B., Fei, Y., Kumar, R. E., He, T., Milsted, D., McDermott, M. J., Gallant, M., Cubuk, E. D., Merchant, A., Kim, H., Jain, A., Bartel, C. J., Persson, K., Zeng, Y., and Ceder, G. An autonomous laboratory for the accelerated synthesis of inorganic materials. *Nature*, 624(7990):86–91, 2023. doi: 10.1038/s41586-023-06734-w.

Tian, M., Gao, L., Zhang, S. D., Chen, X., Fan, C., Guo, X., Haas, R., Ji, P., Krongchon, K., Li, Y., Liu, S., Luo, D., Ma, Y., Tong, H., Trinh, K., Tian, C., Wang, Z., Wu, B., Xiong, Y., Yin, S., Zhu, M., Lieret, K., Lu, Y., Liu, G., Du, Y., Tao, T., Press, O., Callan, J., Huerta, E., and Peng, H. SciCode: A research coding benchmark curated by scientists. *arXiv preprint arXiv:2407.13168*, 2024.

Tom, G., Schmid, S. P., Baird, S. G., Cao, Y., Darvish, K., Hao, H., Lo, S., Pablo-García, S., Rajaonson, E. M., Ramadhan, M., Rodrigues, J., Whitmore, L. S., Antono, E., and Aspuru-Guzik, A. Self-driving laboratories for chemistry and materials science. *Chemical Reviews*, 124(16):9633–9732, 2024. doi: 10.1021/acs.chemrev.4c00055.

Wang, H., Fu, T., Du, Y., Gao, W., Huang, K., Liu, Z., Chandak, P., Liu, S., Van Katwyk, P., Deac, A., et al. Scientific discovery in the age of artificial intelligence. *Nature*, 620:47–60, 2023. doi: 10.1038/s41586-023-06221-2.

Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., and Cao, Y. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2210.03629.

Zhou, S., Xu, F. F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Ou, T., Bisk, Y., Fried, D., Alon, U., and Neubig, G. WebArena: A realistic web environment for building autonomous agents. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2307.13854.

Table 1. Comparison of OPENDISCOVERYTRACE with existing AI science benchmarks. OPENDISCOVERYTRACE is the first benchmark to capture full process traces including error logging, revision tracking, and failure analysis alongside final output evaluation. ✓ = supported, × = not supported, ∼ = partial support.

Benchmark	Process Traces	Error Logging	Revision Tracking	Multi-Domain	Science Specific	Multi-Model
ScienceAgentBench (Chen et al., 2025)	×	×	×	✓	✓	✓
DiscoveryBench (Majumder et al., 2024)	×	×	×	✓	✓	✓
The AI Scientist (Lu et al., 2024)	×	×	×	∼	✓	✓
PaperBench (Starace et al., 2025)	×	×	×	∼	✓	✓
ResearchBench (Liu et al., 2025)	×	×	×	✓	✓	✓
SciCode (Tian et al., 2024)	×	×	×	✓	✓	✓
MLAgentBench (Huang et al., 2024)	∼	∼	×	×	∼	✓
MLE-bench (Chan et al., 2025)	×	×	×	×	×	✓
WebArena (Zhou et al., 2024)	✓	∼	×	×	×	✓
AgentBench (Liu et al., 2024)	✓	∼	×	×	×	✓
OPENDISCOVERYTRACE (ours)	✓	✓	✓	✓	✓	✓

Appendix

A. Full Trace Schema

Listing 1. Complete OPENDISCOVERYTRACE trace schema.

```
{
  "task_id": "string",
  "domain": "drug_discovery | materials_science | genomics | literature",
  "difficulty": "easy | medium | hard",
  "prompt": "string",
  "model": "gpt-5.4 | claude-opus-4-6 | gemini-3.1-pro | ...",
  "trajectory": [
    {
      "step_id": "integer", "timestamp": "ISO8601",
      "phase": "literature_review | hypothesis | experiment_design
             | execution | analysis | revision | conclusion",
      "thought": "string",
      "action": {"type": "string", "tool": "string", "input": {}, "output": {}},
      "observation": "string",
      "error": {"occurred": "boolean", "type": "string", "message": "string"},
      "revision_trigger": "string | null",
      "confidence": "float [0, 1]"
    }
  ],
  "outcome": {
    "success": "boolean", "final_claim": "string",
    "verification": {"method": "string", "result": "string", "score": "float"},
    "failure_type": "string | null", "recovery_attempted": "boolean"
  },
  "metadata": {
    "total_steps": "integer", "total_tokens": "integer",
    "total_tool_calls": "integer", "total_failures": "integer",
    "total_revisions": "integer", "wall_time_seconds": "float"
  }
}
```

B. Benchmark Comparison

C. Pilot Analysis Details

C.1. Dataset Statistics

522 trajectories: 372 frontier (124 each from GPT-5.4, Claude Opus 4.6, Gemini 3.1 Pro) + 120 open-weight (Qwen2.5-7B, Mistral-7B-v0.3, Phi-3.5-mini, Qwen2.5-1.5B; 30 each) + 30 live-retrieval. 99.7% reach conclusions. 31.2% contain errors; 34.4% contain revisions.

Table 2. Inter-model comparison (n=124 per frontier model). *** $p<0.001$, n.s. not significant.

Metric	GPT-5.4	Claude 4.6	Gemini 3.1	Sig.	δ
Steps	3.5±3.5	4.1±4.0	4.8±5.4	n.s.	0.16
Tools	2.5±3.5	2.8±3.2	3.8±5.2	n.s.	0.15
Errors	0.08±0.3	2.5±3.4	0.4±0.8	***	0.61
Revisions	0.4±0.8	1.2±2.0	1.0±2.5	***	0.24
Success	69.0%	69.8%	69.0%	n.s.	—

C.2. Failure Taxonomy

Claude: 462 errors (66.7% tool misuse, 33.1% reasoning, 0.2% hallucination). Gemini: 165 errors (72.1% reasoning, 27.9% tool). GPT-5.4: 61 errors (83.6% reasoning, 16.4% tool).

C.3. Matched-Task Within-Domain Results

Error divergence is significant in every domain ($p<0.0001$); step counts are non-significant in every domain.

C.4. Sequence-Model Baselines

LSTM: Acc=0.693±0.015, AUROC=0.422. Transformer: Acc=0.693±0.015, AUROC=0.369. Both match majority baseline.

C.5. LLM-Based Inter-Annotator Agreement

60-trajectory sample: correctness α=0.629, reasoning α=0.711, tool use α=0.717, autonomy α=0.820.

C.6. Task Examples

Drug Discovery. Easy: "Predict LogP of ibuprofen." Medium: "Find top EGFR kinase inhibitors from ChEMBL, predict ADMET." Hard: "Propose CDK4/6 inhibitor resistance mechanism in TNBC."

Materials Science. Easy: "Band gap of silicon." Medium: "Compare perovskite stability for PV." Hard: "Design HEA synthesis for ductility."

Genomics. Easy: "TP53 protein length." Medium: "KEGG enrichment for HCC genes." Hard: "BRAF V600E across cancer types."

Literature. Easy: "AlphaFold2 CASP14 accuracy." Medium: "CRISPR off-target methods review." Hard: "Gut microbiome-immunotherapy contradiction."